Exercise Session 11 : FLWOR Expressions on Spark.

Agenda :

① Compound index in MongoDB (Session 10)

② Inner workings of the FLOWR expressions
- Tuple streams
- Execution and implementation on Spark.

③ Tips on JSONiq.

MongoDB Compound Indices (session 10)

- Compound keys are "flat" => 1D regardless of # fields
- Non-prefix queries can only use the index as a "filter"

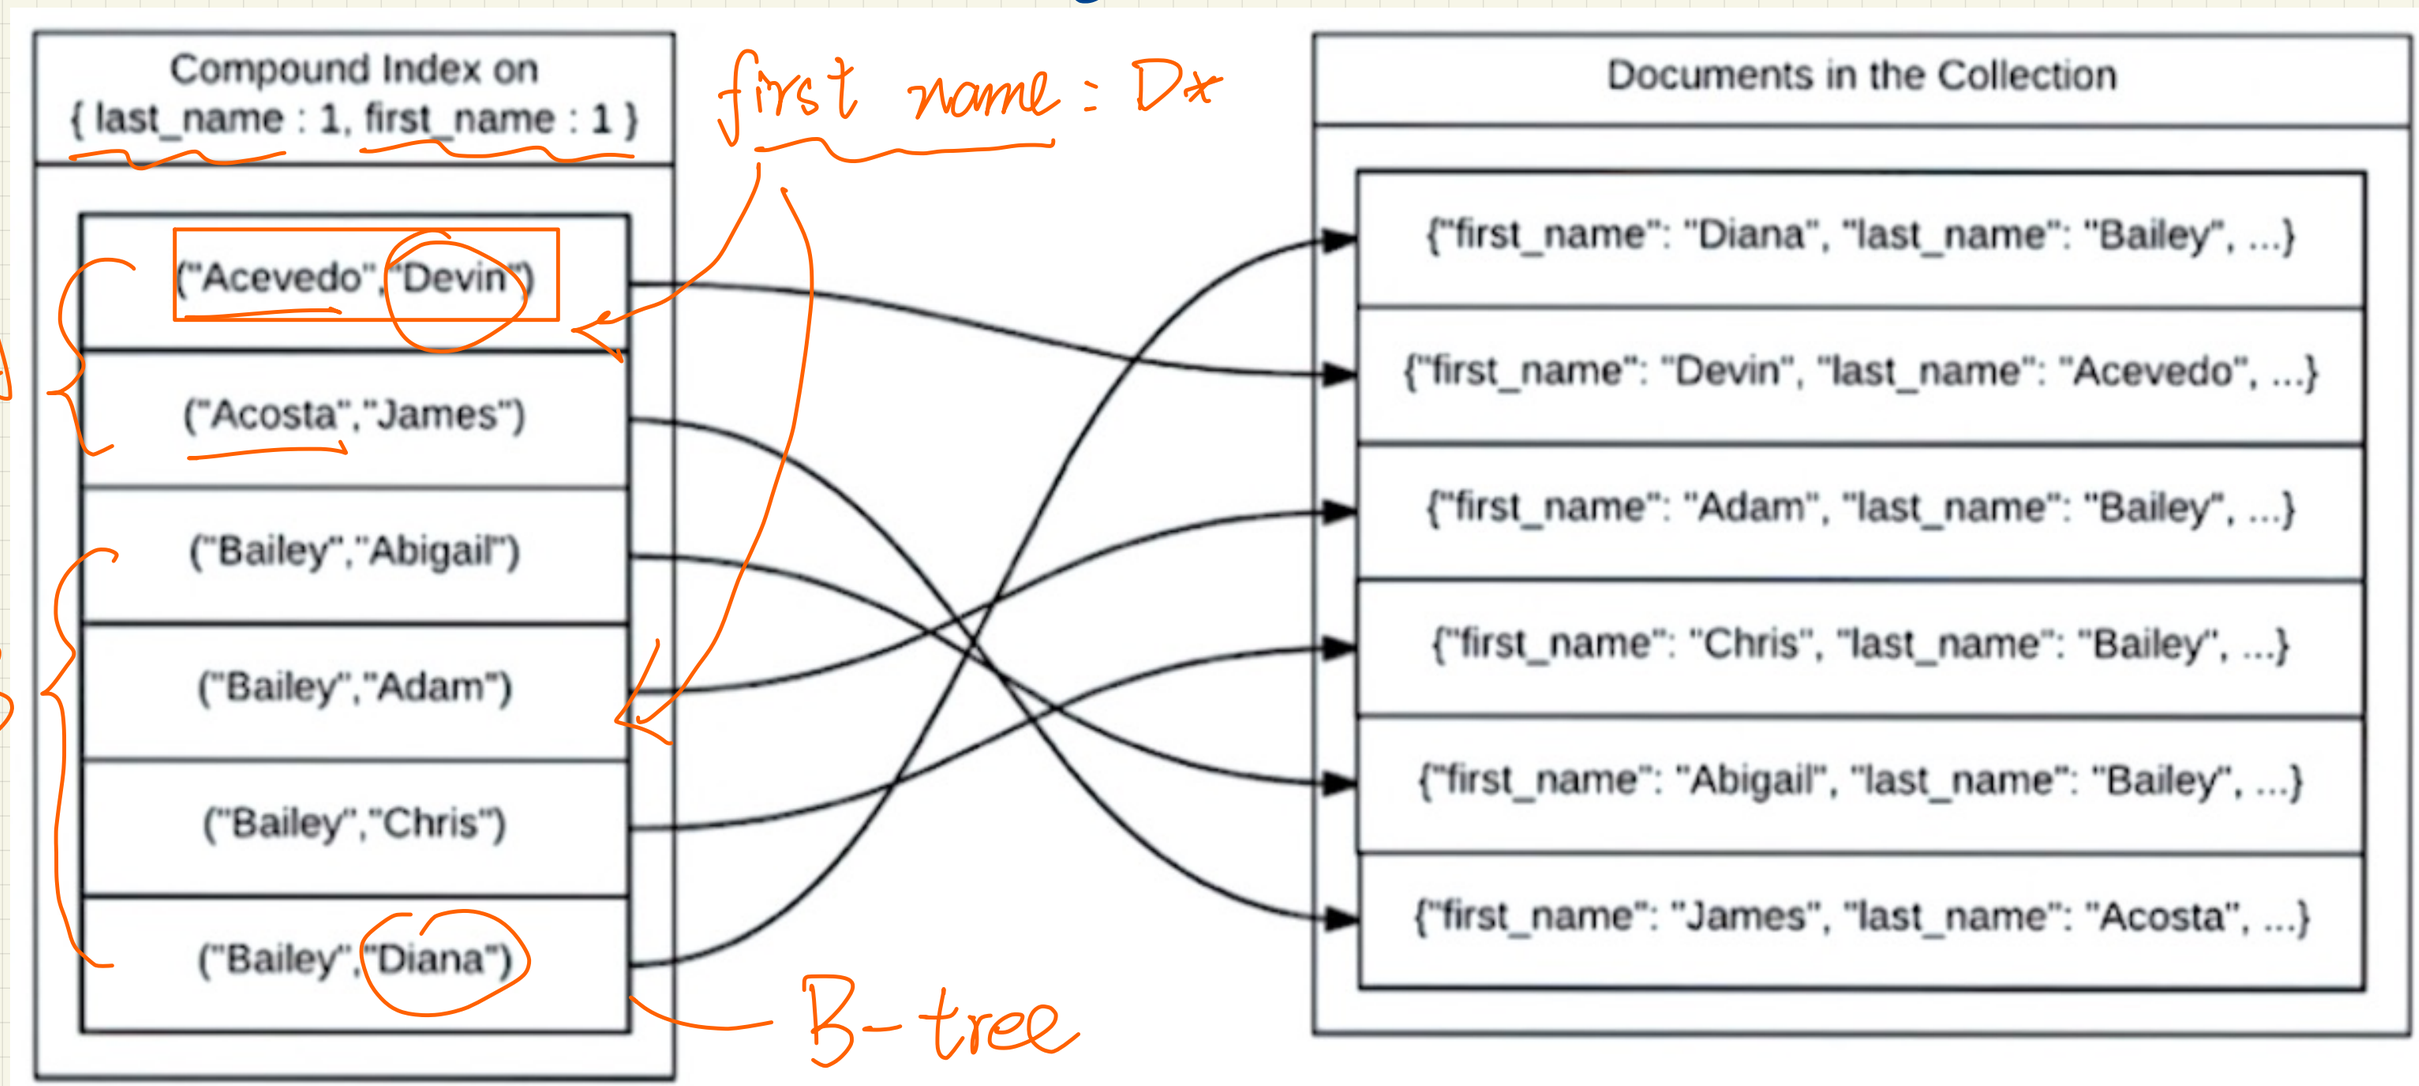

B- tree

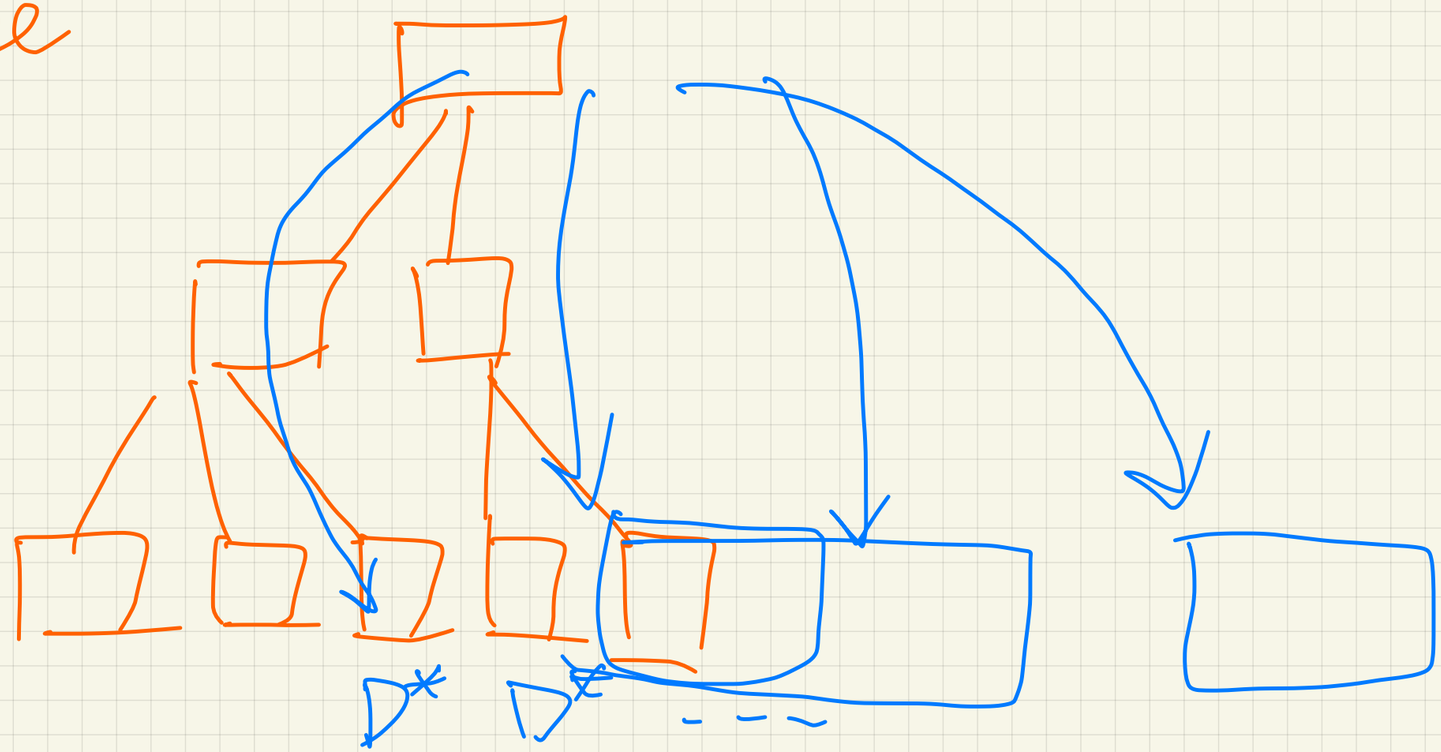

Motivating Example of JSONiq

```
let $arr := ["a", "b", "b", "c", "c", "c"]
```

⟹

```
{"c": 3}
{"a": 1}
{"b": 2}
```

A:

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in $arr[]
group by $x
return {
    $x : count($arr)
}
```

B:

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in distinct-values($arr[])
group by $x
return {
    $x : count($arr[][$$ = $x])
}
```

C:

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in distinct-values($arr[])
group by $x
return {
    $x : count($x)
}
```

D:

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in distinct-values($arr[])
group by $x
return {
    $x : 1
}
```

Outputs :

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in $arr[]
group by $x
return {
    $x : count($arr)
}
```

A ✓

```
Took: 0.0336360931396484 ms
{"c": 3}
{"a": 1}
{"b": 2}
```

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in distinct-values($arr[])
group by $x
return {
    $x : count($arr[][$$ = $x])
}
```

B ✓

```
Took: 0.04476189613342285 ms
{"c": 3}
{"a": 1}
{"b": 2}
```

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c", "c"]
for $x in distinct-values($arr[])
group by $x
return {
    $x : count($x)
}
```

C ✗

```
Took: 0.0374131202697539 ms
{"c": 1}
{"a": 1}
{"b": 1}
```

```
%%jsoniq
let $arr := ["a", "b", "b", "c", "c",
for $x in distinct-values($arr[])
group by $x
return {
    $x : 1
}
```

D ✗

```
Took: 0.03769087791442871 ms
{"c": 1}
{"a": 1}
{"b": 1}
```

FLOWR Semantics Effects

- `let` — Creates a new column
- `for` — Reshapes a list of values into a column
- `where` — Fliters out tuples given a criterion
- `group by` — Collects tuples of the same key to a sequence
- `return` — Emits streams of tuples one by one
  - SELECT

Example Breakdown

A single object

```
let $arr := ["a", "b", "b", "c", "c", "c"]
```

(1)

$arr
["a", "b", "b", "c", "c", "c"]

(2)

```
for $x in $arr[]
```

unpacking → seq. ⇒

["a", "b", "b", "c", "c", "c"]

(3)

```
group by $x
```

$x	$arr
a	[ ]
b	[ ], [ ],
c	[ ], [ ], [ ]

seq.

tuple

$x	$arr
a	["a", "b", "b", "c", "c", "c"]
b	["a", "b", "b", "c", "c", "c"]
b	["a", "b", "b", "c", "c", "c"]
c	["a", "b", "b", "c", "c", "c"]
c	["a", "b", "b", "c", "c", "c"]
c	["a", "b", "b", "c", "c", "c"]

aggregation on each tuple

(4)

```
return {
    $x : count($arr)
}
```

More Examples from Lectures

```
for $x in 1 to 5
let $y := $x - 2
return $x * $y
```

x
1
2
3
4
5

x	y
1	-1
2	0
3	1
4	2
5	3

-1
0
3
8
15

```
for $x in 1 to 5
group by $y := $x mod 2
return {
  "foo" : $y,
  "bar" : count($x)
}
```

x
1
2
3
4
5

x	y
1	1
2	0
3	1
4	0
5	1

foo

x	y
(1, 3, 5)	1
(2, 4)	0

a sequence

```
{"foo":1,"bar":3}
{"foo":0,"bar":2}
```

```
for $x in 1 to 10
where $x - 2 gt 5
return $x * $x
```

x
1
2
3
4
5
6
7
8
9
10

(where)
false
false
false
false
false
false
false
true
true
true

output stream

x
8
9
10

64
81
100

Execution

Streaming

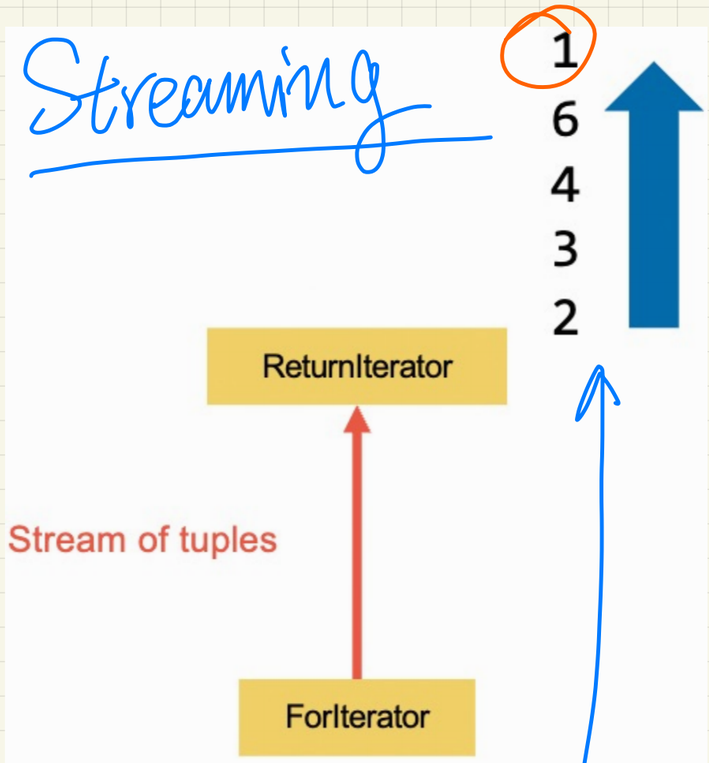


```
$i: { "field" : 1 }
$i: { "field" : 6 }
$i: { "field" : 4 }
$i: { "field" : 3 }
$i: { "field" : 2 }
```

Parallel

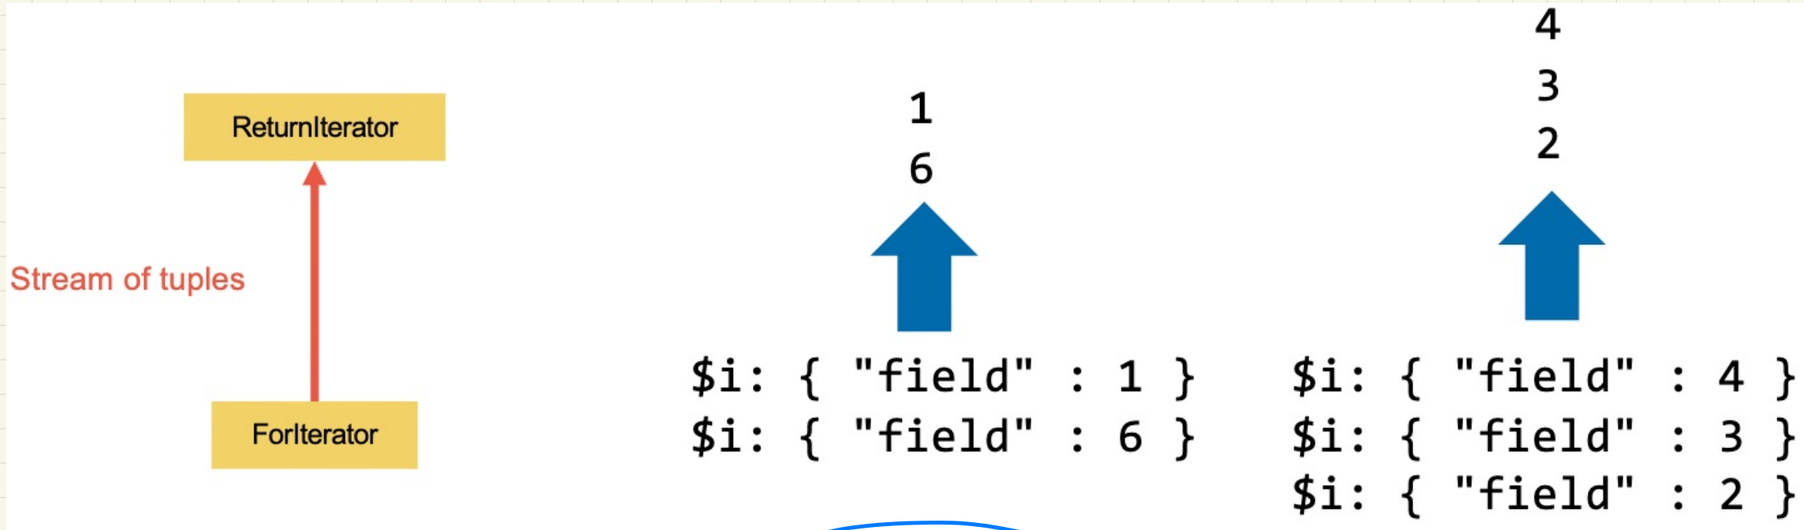


```
$i: { "field" : 1 }
$i: { "field" : 6 }
$i: { "field" : 4 }
$i: { "field" : 3 }
$i: { "field" : 2 }
```

RDD

Implimentation atop Spark

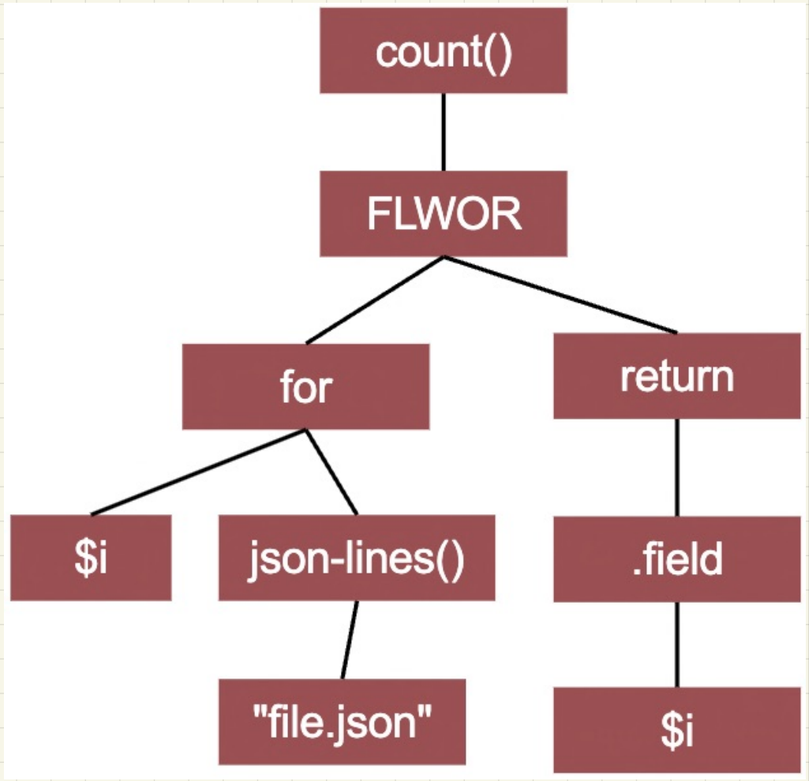

AST

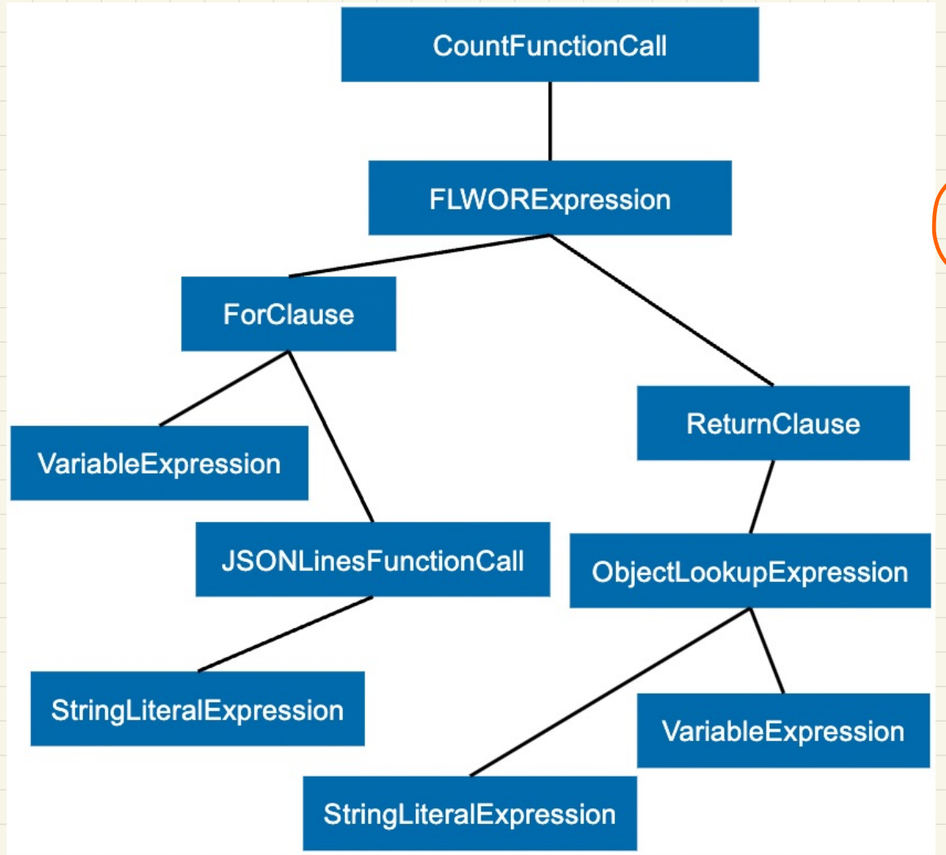

Expression Tree

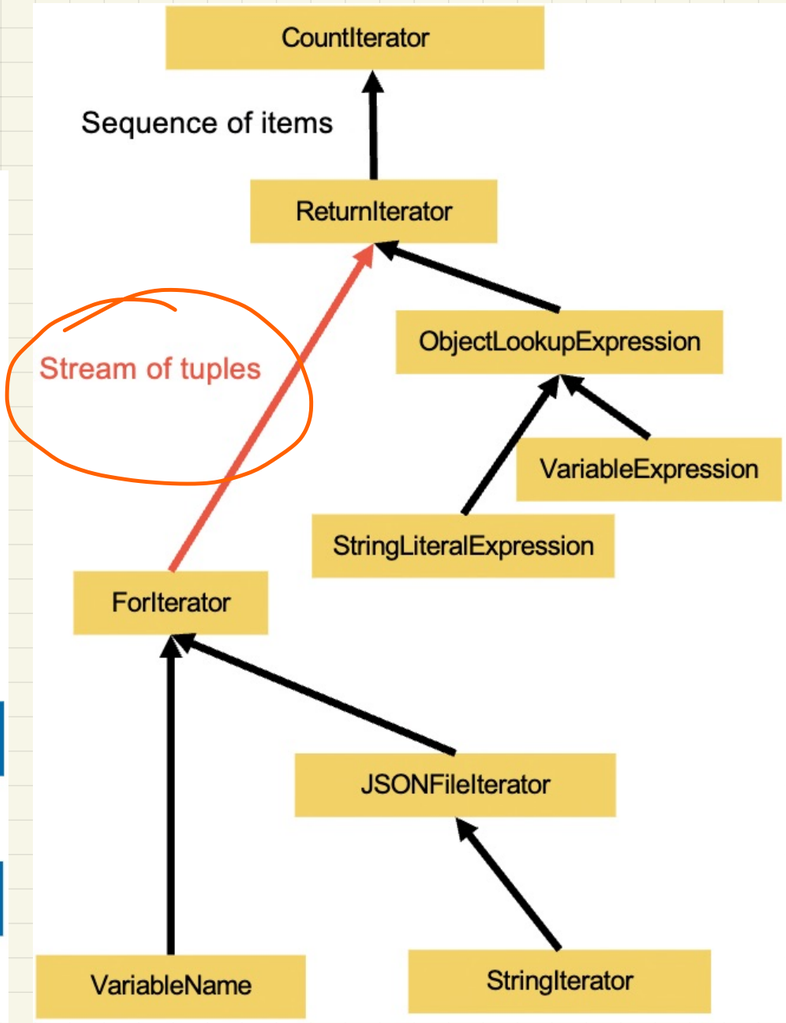

Iterator Tree

Spark Implementation

new coC

```
for $person in json-file("persons")
...
```

Implementation:
mapping of the RDD to a single-column DataFrame

serialized binary format

person
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }
{ "Name" : { "First" : "Srinivasa", "Last" : "Ramanujan" }, "Countries" : ["IN", "UK"] }
{ "Name" : { "First" : "Kurt", "Last" : "Gödel" }, "Countries" : ["CZ", "A", "US"] }
{ "Name" : { "First" : "John", "Last" : "Nash" }, "Countries" : "US" }
{ "Name" : { "First" : "Alan", "Last" : "Turing" }, "Countries" : "UK" }
{ "Name" : { "First" : "Maryam", "Last" : "Mirzakhani" }, "Countries" : ["IR", "US"] }
{ "Name" : "Pythagoras", "Countries" : ["GR"] }
{ "Name" : { "First" : "Nicolas", "Last" : "Bourbaki" }, "Number" : 9, "Countries" : null }

```
for $person in json-file("persons")
for $country in flatten($person.Countries[])[$$ ne null]
...
```

let

Cartesian Product

Implementation:
Spark SQL + UDF

person	country
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	D
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	I
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	CH
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	A
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	BE
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	US
{ "Name" : { "First" : "Srinivasa", "Last" : "Ramanujan" }, "Countries" : ["IN", "UK"] }	IN
{ "Name" : { "First" : "Srinivasa", "Last" : "Ramanujan" }, "Countries" : ["IN", "UK"] }	UK
{ "Name" : { "First" : "Kurt", "Last" : "Gödel" }, "Countries" : ["CZ", "A", "US"] }	CZ
{ "Name" : { "First" : "Kurt", "Last" : "Gödel" }, "Countries" : ["CZ", "A", "US"] }	A

```
for $person in json-file("persons")
for $country in flatten($person.Countries[])[$$ ne null]
where $country eq "US"
...
```

Implementation:
Spark SQL + UDF

person	country
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	US
{ "Name" : { "First" : "Kurt", "Last" : "Gödel" }, "Countries" : ["CZ", "A", "US"] }	US
{ "Name" : { "First" : "John", "Last" : "Nash" }, "Countries" : "US" }	US
{ "Name" : { "First" : "Maryam", "Last" : "Mirzakhani" }, "Countries" : ["IR", "US"] }	US

```
for $person in json-file("persons")
for $country in flatten($person.Countries[])[$$ ne null]
where $country eq "US"
let $first-name eq $person.Name.First
...
```

Implementation:
Spark SQL + UDF

person	country	first-name
{ "Name" : { "First" : "Albert", "Last" : "Einstein" }, "Countries" : ["D", "I", "CH", "A", "BE", "US"] }	US	Albert
{ "Name" : { "First" : "Kurt", "Last" : "Gödel" }, "Countries" : ["CZ", "A", "US"] }	US	Kurt
{ "Name" : { "First" : "John", "Last" : "Nash" }, "Countries" : "US" }	US	John
{ "Name" : { "First" : "Maryam", "Last" : "Mirzakhani" }, "Countries" : ["IR", "US"] }	US	Maryam

```
for $person in json-file("persons")
for $country in flatten($person.Countries[])[$$ ne null]
where $country eq "US"
let $first-name eq $person.Name.First
return $first-name
```

emit
one by one

Implementation:
Mapping back to an RDD

"Albert"
"Kurt"
"John"
"Maryam"

Final Tips

① Iterate over arrays:

```
1 %%jsoniq
2 for $i in [1, 2, 3]
3 return $i
```

```
Took: 0.03023815155029297 ms
[1, 2, 3]
```

→ a single item

```
1 %%jsoniq
2 for $i in [1, 2, 3][]
3 return $i
```

3 items (ints)

```
Took: 0.03351902961730957 ms
1
2
3
```

② Implicit existential quantification:

```
1 %%jsoniq
2 10 eq 5
```

```
Took: 0.029591083526611328 ms
false
```

```
1 %%jsoniq
2 1 to 10 = 5
```

```
Took: 0.04393911361694336 ms
true
```

[1, 10] [9, 20]

```
1 %%jsoniq
2 1 to 10 > 9 to 20
```

```
Took: 0.03199887275695801 ms
true
```

```
1 %%jsoniq
2 1 to 10 > 11 to 20
```

```
Took: 0.04635882377624512 ms
false
```